

Appendix A: Code

run.sh

```
out_dir="results"
trials=10
# Must be powers of two.
# I go from 1M (131072 float2's) to 512M (33554432 float2's)
arr_size=(131072 524288 2097152 8388608 33554432)
# Must also be powers of two. I pass this twice since I use num_blks=num_thds.
num_thds=(2 4 8 16 32)
# The names of the executables
seq_file="./sequential-fft"
par_file="./parallel-fft"
# Clean the result directory
rm -rf $out_dir
mkdir $out_dir
# Run sequential
for size in "${arr_size[@]}"
do
    echo 1 $size
    out=$(seq_file $size $trials)
    echo -e $size'\t'$out >> $out_dir/1.txt
done
# Run parallel
for threads in "${num_thds[@]}"
do
    # Total number of threads = num_blk * num_thd
    exprans='expr $threads \/* $threads'
    for size in "${arr_size[@]}"
    do
        echo $exprans $size
        out=$(par_file $size $trials $threads $threads)
        echo -e $size'\t'$out >> $out_dir/$exprans.txt
    done
done
```

sequential-fft.cu

```
// This program takes n and trial count as parameters and nothing more.
// It is assumed that n is a power of 2.
// Compiles with nvcc -std=c++11 -rdc=true -arch=compute_50 -code=sm_50
#include <chrono>
#include <cstdlib>
#include <ctime>
#include <iostream>
#include <utility>

#include <math_functions.h>

#include "include/fast-fourier.h"

using namespace std;
using namespace chrono;
using namespace fast_fourier;

void gen_array(cfloat* output, int n);
long double average(long double* in, int n);
long double std_dev(long double* in, int n, long double average);

__global__
void run_test(cfloat* input, cfloat* output, int n, bool* binary_stor)
{
    fast_fourier_transform(input, output, n, binary_stor);
}

int main(int argc, char** argv)
{
    if (argc < 3)
    {
        cerr << "Usage is " << argv[0] << " n num_trials" << endl;
        return 1;
    }

    int n(atoi(argv[1]));
    int trial_count(atoi(argv[2]));

    cfloat* input(new cfloat[n]);
    cfloat* output(new cfloat[n]);
    cfloat* d_input(nullptr);
    cfloat* d_output(nullptr);
    bool* binary_stor(nullptr);

    long double times[trial_count];
    high_resolution_clock::time_point tp2, tp1;
    duration<long double, ratio<1,1000> > time_span;

    // Allocate device arrays
    if (cudaMalloc( &d_input, sizeof(cfloat) * n ) != cudaSuccess)
    {
        auto t = cudaGetLastError();
        cout << "Failed to allocate input: "
              << cudaGetErrorName(t) << ", "
              << cudaGetErrorString(t) << endl;
        return 1;
    }
    if (cudaMalloc( &d_output, sizeof(cfloat) * n ) != cudaSuccess)
    {
        auto t = cudaGetLastError();
        cout << "Failed to allocate output: "
              << cudaGetErrorName(t) << ", "
              << cudaGetErrorString(t) << endl;
        return 1;
    }
    if (cudaMalloc( &binary_stor, sizeof(bool) * ilogbf(n) ) != cudaSuccess)
    {
        auto t = cudaGetLastError();
        cout << "Failed to allocate boolean storage: "
              << cudaGetErrorName(t) << ", "
              << cudaGetErrorString(t) << endl;
        return 1;
    }

    // Run experiment
```

```

for (int j(0) ; j < trial_count ; j++)
{
    // Generate random input
    gen_array(input, n);

    // Run the test
    tp1 = system_clock::now();
    // Copy the input array to the GPU
    if (cudaMemcpy( d_input, input, (long) n * sizeof(cfloat), cudaMemcpyHostToDevice ) != cudaSuccess)
    {
        auto t = cudaGetLastError();
        cout << "Iteration: " << j
              << " Input failed to copy: "
              << cudaGetErrorName(t) << ", "
              << cudaGetErrorString(t) << endl;
        return 1;
    }
    run_test<<<1,1>>>(d_input, d_output, n, binary_stor);
    if (cudaMemcpy( output, d_output, (long) n * sizeof(cfloat), cudaMemcpyDeviceToHost ) != cudaSuccess)
    {
        auto t = cudaGetLastError();
        cout << "Iteration: " << j
              << " Output failed to copy: "
              << cudaGetErrorName(t) << ", "
              << cudaGetErrorString(t) << endl;
        return 1;
    }
    tp2 = system_clock::now();

    time_span    = duration_cast< duration<long double, ratio<1,1000> > >(tp2 - tp1);
    times[j]     = time_span.count();
}

// Calculate statistics
long double av(average(times, trial_count));
long double sd(std_dev(times, trial_count, av));

cout << av << "\t" << sd << endl;

cudaFree( binary_stor );
cudaFree( d_input );
cudaFree( d_output );
return 0;
}

void gen_array(cfloat* output, int n)
{
    srand(time(nullptr));

    for (int j = 0; j < n; j++)
        output[j] = cfloat(rand(), rand());
}

long double    average(long double* in, int n)
{
    long double s(0.0);

    for (int j(0) ; j < n ; j++)
        s += in[j];

    return s/n;
}

long double    std_dev(long double* in, int n, long double average)
{
    long double var = 0;
    long double tmp = 0;

    for (int i = 0 ; i < n ; i++)
    {
        tmp = (in[i] - average);
        var += tmp * tmp;
    }

    long double stdDev = sqrt(var/n);

    return stdDev;
}

```


parallel-fft.cu

```
// This program takes n, block count, thread count, and number of trials as parameters.
// It is assumed that these are all powers of 2.
// Compiles with nvcc -std=c++11 -rdc=true -arch=compute_50 -code=sm_50
#include <chrono>
#include <cstdlib>
#include <ctime>
#include <iostream>
#include <utility>

#include "include/fast-fourier.h"

using namespace std;
using namespace chrono;
using namespace fast_fourier;

void gen_array(cfloat* output, int n);
long double average(long double* in, int n);
long double std_dev(long double* in, int n, long double average);

int main(int argc, char** argv)
{
    if (argc < 3)
    {
        cerr << "Usage is " << argv[0] << " n num_trials num_blk num_thd" << endl;
        return 1;
    }

    int n(atoi(argv[1]));
    int trial_count(atoi(argv[2]));
    int num_blk(atoi(argv[3]));
    int num_thd(atoi(argv[4]));

    cfloat* input(new cfloat[n]);
    cfloat* output(new cfloat[n]);
    cfloat* d_input(nullptr);
    cfloat* d_output(nullptr);
    bool* binary_stor(nullptr);

    long double times[trial_count];
    high_resolution_clock::time_point tp2, tp1;
    duration<long double, ratio<1,1000> > time_span;

    // Allocate device arrays
    if (cudaMalloc( &d_input, sizeof(cfloat) * n ) != cudaSuccess)
    {
        auto t = cudaGetLastError();
        cout << "Failed to allocate input: "
              << cudaGetErrorName(t) << ", "
              << cudaGetErrorString(t) << endl;
        return 1;
    }
    if (cudaMalloc( &d_output, sizeof(cfloat) * n ) != cudaSuccess)
    {
        auto t = cudaGetLastError();
        cout << "Failed to allocate output: "
              << cudaGetErrorName(t) << ", "
              << cudaGetErrorString(t) << endl;
        return 1;
    }
    if (cudaMalloc( &binary_stor, sizeof(bool) * ilogbf(n) * num_thd * num_blk) != cudaSuccess)
    {
        auto t = cudaGetLastError();
        cout << "Failed to allocate boolean storage: "
              << cudaGetErrorName(t) << ", "
              << cudaGetErrorString(t) << endl;
        return 1;
    }

    // Run experiment
    for (int j(0) ; j < trial_count ; j++)
    {
        // Generate random input
        gen_array(input, n);

        // Run the test
```

```

    tp1 = system_clock::now();
    // Copy the input array to the GPU
    if (cudaMemcpy( d_input, input, (long) n * sizeof(cfloat), cudaMemcpyHostToDevice ) != cudaSuccess)
    {
        auto t = cudaGetLastError();
        cout << "Iteration: " << j
              << " Input failed to copy: "
              << cudaGetErrorName(t) << ", "
              << cudaGetErrorString(t) << endl;
        return 1;
    }
    fast_fourier_transform<<<1,1>>>(d_input, d_output, n, num_blk, num_thd, binary_stor);
    if (cudaMemcpy( output, d_output, (long) n * sizeof(cfloat), cudaMemcpyDeviceToHost ) != cudaSuccess)
    {
        auto t = cudaGetLastError();
        cout << "Iteration: " << j
              << " Output failed to copy: "
              << cudaGetErrorName(t) << ", "
              << cudaGetErrorString(t) << endl;
        return 1;
    }
    tp2 = system_clock::now();

    time_span    = duration_cast< duration<long double, ratio<1,1000> > >(tp2 - tp1);
    times[j]     = time_span.count();
}

// Calculate statistics
long double av(average(times, trial_count));
long double sd(std_dev(times, trial_count, av));

cout << av << "\t" << sd << endl;

cudaFree( binary_stor );
cudaFree( d_input );
cudaFree( d_output );
return 0;
}

void gen_array(cfloat* output, int n)
{
    srand(time(nullptr));

    for (int j = 0; j < n; j++)
        output[j] = cfloat(rand(), rand());
}

long double average(long double* in, int n)
{
    long double s(0.0);

    for (int j(0) ; j < n ; j++)
        s += in[j];

    return s/n;
}

long double std_dev(long double* in, int n, long double average)
{
    long double var = 0;
    long double tmp = 0;

    for (int i = 0 ; i < n ; i++)
    {
        tmp = (in[i] - average);
        var += tmp * tmp;
    }

    long double stdDev = sqrt(var/n);

    return stdDev;
}

```

```

include/fast-fourier.h

#ifndef FAST_FOURIER_H
#define FAST_FOURIER_H

#include <thrust/complex.h>

namespace fast_fourier
{
    typedef thrust::complex<float> cfloat;

    /// Does fast fourier transform.
    /// @param x The input vector (MUTABLE!!!).
    /// @param y The output vector.
    /// @param n The size of the vectors.
    /// @param binary_stor A place to store binary representations.
    __host__ __device__
    void fast_fourier_transform(cfloat* x, cfloat* y, unsigned n, bool* binary_stor);
    /// Does parallel fast fourier transform.
    /// @param x The input vector (MUTABLE!!!).
    /// @param y The output vector.
    /// @param n The size of the vectors.
    /// @param blk_count The number of blocks you wish to spawn.
    /// @param thd_count The number of threads per block.
    /// @param binary_stor A place to store binary representations.
    __global__
    void fast_fourier_transform(cfloat* x, cfloat* y, unsigned n, int blk_count, int thd_count, bool* binary_stor);
    /// Does discrete fourier transform.
    /// @param x The input vector.
    /// @param n The size of the vector.
    /// @return A pointer to the output vector.
    __host__ __device__
    cfloat* discrete_fourier_transform(cfloat* x, unsigned n);
}

#endif

```

```

include/src/fast-fourier.cu

// Compiles with nvcc -std=c++11 -rdc=true -arch=compute_50 -code=sm_50
#include <thrust/fill.h>
#include <thrust/copy.h>
#include <math_constants.h>
#include <math_functions.h>

#include "../fast-fourier.h"

using thrust::copy_n;
using thrust::exp;
using thrust::fill_n;

using fast_fourier::cfloat;

/// Increments a number represented in binary.
/// @param i Our binary representation.
/// @param lg_n The size of the bool vector (it had better be lg n!).
__host__ __device__
void binary_inc(bool* i, int lg_n);
/// Converts a binary representation to an integer.
/// @param i Our binary representation.
/// @param lg_n The size of the bool vector (it had better be lg n!).
/// @return The decimal value
__host__ __device__
int bin2dec(const bool* i, int lg_n);
/// Writes a binary representation to a bit vector.
/// @param i Our binary representation.
/// @param l The int to convert.
/// @param lg_n The size of the bool vector (it had better be lg n!).
__host__ __device__
void dec2bin(bool* i, int l, int lg_n);
/// Our wierd binary exponent.
/// @param l Our binary representation.
/// @param lg_n The size of the bool vector (it had better be lg n!).
/// @param m The current iteration.
__host__ __device__
int wierd_bin_thingy(const bool* l, int lg_n, int m);
/// Returns the nth root of unity raised to the k
__host__ __device__
cfloat k_root_unity(int k, int n);

/// The inner loop of FFT
/// @param r The vector to which we write.
/// @param s The vector from which we read.
/// @param lg_n lg n.
/// @param blk_off The block offset (used for reading from binary_stor).
/// @param thd_off The thread offset (used for reading from binary_stor).
/// @param n The size of our I/O vectors.
/// @param m The current iteration.
/// @param binary_stor our bit vector where we store binary representations.
/// @param thd_count The number of threads per block.
__global__
void transformer(cfloat* r, cfloat* s, unsigned lg_n, unsigned blk_off,
                unsigned thd_off, unsigned n, int m, bool* binary_stor, int thd_count);
/// A parallel copy algorithm.
/// @param src The source.
/// @param dst The destination.
/// @param n The size of the vectors.
/// @param blk_off The number of things each block gets.
/// @param thd_off The number of things each thread gets.
__global__
void parallel_copy(const cfloat* src, cfloat* dst, unsigned n, unsigned blk_off,
                 unsigned thd_off);

cfloat* fast_fourier::discrete_fourier_transform(cfloat* x, unsigned n)
{
    cfloat* y(new cfloat[n]);

    fill_n(y, n, cfloat(0.0f));

    for (int j(0) ; j < n ; j++)
        for (int k(0) ; k < n ; k++)
            y[j] += x[k] * k_root_unity(k*j, n);

    return y;
}

```



```

}

void fast_fourier::fast_fourier_transform(cfloat* x, cfloat* y, unsigned n,
                                         bool* binary_stor)
{
    cfloat* s(x);
    cfloat* r(y);
    cfloat* tmp_ptr;
    int lg_n(ilogbf(n));
    int j,k,u_exp;
    bool* l_bi(binary_stor);
    bool tmp(false);

    for (int j(0) ; j < lg_n ; j++ )
        l_bi[j] = false;

    for (int m(0) ; m < lg_n ; m++)
    {
        tmp_ptr = s;
        s = r;
        r = tmp_ptr;

        for (int l(0) ; l < n ; l++)
        {
            tmp = l_bi[l];
            l_bi[l] = false;

            j = bin2dec(l_bi, lg_n);
            k = j + (int)exp2f(lg_n - m - 1);

            l_bi[l] = tmp;

            u_exp = wierd_bin_thingy(l_bi, lg_n, m);
            r[l] = s[j] + s[k] * k_root_unity(u_exp, n);

            binary_inc(l_bi, lg_n);
        }
    }

    for (int j(0) ; j < n ; j++)
        y[j] = r[j];
}

__global__
void fast_fourier::fast_fourier_transform(cfloat* x, cfloat* y, unsigned n,
                                         int blk_count, int thd_count, bool* binary_stor)
{
    int lg_n(ilogbf(n));

    int blk_off = n / blk_count;
    int thd_off = blk_off / thd_count;

    cfloat *r(x), *s(y);
    cfloat* tmp_ptr;

    for (int m(0) ; m < lg_n ; m++)
    {
        // Swap s and r so the last output becomes the new input
        tmp_ptr = r;
        r = s;
        s = tmp_ptr;

        // Perform the next step of the transform
        transformer<<<blk_count, thd_count>>>(r, s, lg_n, blk_off, thd_off, n, m, binary_stor, thd_count);
        cudaDeviceSynchronize();
    }

    // Copy r into y
    parallel_copy<<<blk_count, thd_count>>>(r, y, n, blk_off, thd_off);
    cudaDeviceSynchronize();

    // delete[] r, s;
}

__global__
void transformer(cfloat* r, cfloat* s, unsigned lg_n, unsigned blk_off,
                unsigned thd_off, unsigned n, int m, bool* binary_stor, int thd_count)

```

```

{
    int      l_min( blockIdx.x * blk_off + threadIdx.x * thd_off );
    int      l_max( l_min + thd_off );
    int      j, k, u_exp;

    // bool*    l_bi(new bool[lg_n]);
    bool*    l_bi(binary_stor + (blockIdx.x * thd_count + threadIdx.x) * lg_n);
    bool      tmp;

    dec2bin(l_bi, l_min, lg_n);

    for (int l(l_min) ; l < l_max ; l++)
    {
        tmp = l_bi[l];

        l_bi[l]    = false;
        j          = bin2dec(l_bi, lg_n);
        k          = j + (int)exp2f(lg_n - m - 1);
        l_bi[l]    = tmp;

        u_exp      = wierd_bin_thingy(l_bi, lg_n, m);
        r[l]       = s[j] + s[k] * k_root_unity(u_exp, n);

        binary_inc(l_bi, lg_n);
    }

    // delete[] l_bi;
}

__global__
void parallel_copy(const cfloat* src, cfloat* dst, unsigned n, unsigned blk_off,
                  unsigned thd_off)
{
    int      l_min( blockIdx.x * blk_off + threadIdx.x * thd_off );
    int      l_max( l_min + thd_off );

    for (int l(l_min) ; l < l_max ; l++)
        dst[l] = src[l];
}

void binary_inc(bool* i, int lg_n)
{
    bool flag(true);

    for (int j(lg_n - 1) ; j > -1 && flag ; j--)
    {
        flag = i[j];
        i[j] = !flag;
    }
}

int bin2dec(const bool* i, int lg_n)
{
    int m(0), two_j(1);

    for (int j(lg_n - 1) ; j > -1 ; j--)
    {
        m += (int)i[j] * two_j;
        two_j *= 2;
    }

    return m;
}

int wierd_bin_thingy(const bool* l, int lg_n, int m)
{
    int exponent(0), two_j(1);

    for (int j(0) ; j < m + 1 ; j++)
    {
        exponent += l[j] * two_j;
        two_j *= 2;
    }

    for (int j(m + 1) ; j < lg_n ; j++)
        exponent *= 2;
}

```

```

    return exponent;
}

cfloat k_root_unity(int k, int n)
{
    float e_img(2.0f * ((float)k) * CUDART_PI_F / ((float) n));
    cfloat exponent(cfloat(0,e_img));

    return exp(exponent);
}

void dec2bin(bool* i, int l, int lg_n)
{
    for (int j(lg_n - 1) ; j > -1 ; j--)
    {
        i[j] = (bool) (l % 2);
        l /= 2;
    }
}

```

Appendix B: Results

1.txt

131072	7373.71	9.24972
524288	35790.6	6.01608
2097152	162529	20.0583
8388608	773815	44.6278
33554432	3.42766e+06	295.368

4.txt

131072	1747.24	9.69852
524288	8465.2	8.9587
2097152	38511.3	9.40846
8388608	183374	18.6204
33554432	813669	64.7037

16.txt

131072	445.596	9.95591
524288	2145.66	9.62443
2097152	9747.84	7.09251
8388608	46412	12.7099
33554432	205829	6.82086

64.txt

131072	117.242	5.8156
524288	552.491	9.52933
2097152	2508.03	7.97299
8388608	11884	8.44934
33554432	52810.5	7.01352

256.txt

131072	33.8304	1.67292
524288	148.929	6.68021
2097152	668.6	7.76817
8388608	3142.33	7.90265
33554432	14307.3	10.0786

1024.txt

131072	11.0459	0.0331117
524288	46.11	2.78446
2097152	201.333	6.6144
8388608	975.441	8.89214
33554432	4653.17	6.70649