

# GPU Accelerated Fast Fourier Transform

Matt McCarthy

*Christopher Newport University*  
*matthew.mccarthy.12@cnu.edu*

December 2015

**Abstract** We empirically investigate the performance benefits of parallel fast Fourier transform running on the GPU over a sequential version running on the GPU.

## 1 Background

### 1.1 Discrete Fourier Transform

The discrete Fourier transform is a mathematical transformation that takes a set of Complex-valued signals and outputs a set of Complex-valued frequencies. For an  $n$ -dimensional Complex-valued vector  $\mathbf{X}$ , the discrete Fourier transform  $\mathbf{Y} = \mathcal{F}(X)$  is given by

$$Y_j = \sum_{k=0}^n x_k \omega^{jk}$$

where  $\omega$  is the  $n$ -th root of unity,  $e^{2\pi i/n}$ . Since  $\mathbf{Y}$  is an  $n$ -dimensional, Complex-valued vector, we can see that the discrete Fourier transform has a complexity of  $\Theta(n^2)$ .

### 1.2 Fast Fourier Transform

Furthermore, we can split the discrete Fourier transform into even and odd sums for  $n = 2m$ , yielding

$$Y_j = \sum_{k=0}^m x_{2k} \omega^{2jk} + \omega^j \sum_{k=0}^m x_{2k+1} \omega^{2jk}$$

which is two separate discrete Fourier transforms. Suppose  $n = 2^k$ . If we iterate this process, we get the following algorithm called the one-dimensional, unordered radix 2, fast Fourier transform.

```
1: function R-FFT( $\mathbf{X}, \mathbf{Y}, n, \omega$ )  
2:   if  $n=1$  then
```

```
3:      $y_0 = x_0$   
4:   else  
5:     Let  $\mathbf{Q} = \mathbf{0}, \mathbf{T} = \mathbf{0} \in \mathbb{C}^n$   
6:     Let  $\mathbf{X}_e = (x_0, x_2, \dots, x_{n-2})$   
7:     Let  $\mathbf{X}_o = (x_1, x_3, \dots, x_{n-1})$   
8:     R-FFT( $\mathbf{X}_e, \mathbf{Q}_e, n/2, \omega^2$ )  
9:     R-FFT( $\mathbf{X}_o, \mathbf{T}_o, n/2, \omega^2$ )  
10:    for all  $j \in \{0, 1, \dots, n-1\}$  do  
11:       $y_j = q_j \bmod n/2 + \omega^j t_j \bmod n/2$   
12:    end for  
13:  end if  
14: end function
```

#### 1.2.1 Cooley Tukey

Furthermore, we have an iterative formulation of the prior algorithm, called the Cooley Tukey algorithm for one-dimensional, unordered radix 2, fast Fourier transforms.

```
1: function I-FFT( $\mathbf{X}, \mathbf{Y}, n$ )  
2:    $t := \lg n$   
3:    $\mathbf{R} = \mathbf{X}$   
4:   for  $m = 0$  to  $t-1$  do  
5:      $\mathbf{S} = \mathbf{R}$   
6:     for  $l = 0$  to  $n-1$  do  
7:       Let  $(b_0 b_1 \dots b_{t-1})$  be the binary expansion of  $l$   
8:        $j := (b_0 \dots b_{m-1} 0 b_{m+1} \dots b_{t-1})$   
9:        $k := (b_0 \dots b_{m-1} 1 b_{m+1} \dots b_{t-1})$   
10:       $r_i := s_j + s_k \omega^{(b_m b_{m-1} \dots b_0 0 \dots 0)}$   
11:    end for  
12:  end for  
13:   $\mathbf{Y} := \mathbf{R}$   
14: end function
```

### 1.3 Parallelization

For our parallelization, we use a simplified version of the binary exchange algorithm, a parallelization of the Cooley Tukey algorithm designed for use on a hypercube. Since our implementation runs on a single GPU, any thread can access any memory location via a pointer. However, this also complicates the matter by introducing a potential for data races. We solve this by modifying the algorithm to work as follows.

```
1: function PAR-FFT(X,Y, $n$ )
2:    $t := \lg n$ ,  $BLK := n/p$ 
3:   R = X
4:   S = 0
5:   for  $m = 0$  to  $t - 1$  do
6:     Swap pointers R and S
7:     spawn process for  $l = 0$  to  $BLK - 1$  do
8:       for  $c = l \cdot BLK$ , to  $l \cdot (BLK + 1)$  do
9:         Let  $(b_0 b_1 \dots b_{t-1})$  be the binary ex-
pansion of  $c$ 
10:         $j := (b_0 \dots b_{m-1} 0 b_{m+1} \dots b_{t-1})$ 
11:         $k := (b_0 \dots b_{m-1} 1 b_{m+1} \dots b_{t-1})$ 
12:         $r_i := s_j + s_k \omega^{(b_m b_{m-1} \dots b_0 0 \dots 0)}$ 
13:      end for
14:    end spawn
15:    sync
16:  end for
17:  Y := R
18: end function
```

## 2 Experimental Design

### 3 Test Environment

#### 3.1 Test System

#### 3.2 Test Program

## 4 Results

#### 4.1 Linear Speedup

## 5 Conclusion

## Appendix